



## Chapter 6: Basic SQL

---

CS-6360 Database Design

Instructor: Chris Irwin Davis

[chrisirwindavis@utdallas.edu](mailto:chrisirwindavis@utdallas.edu)

- 
- 6.1 – SQL Data Definition and Data Types
  - 6.2 – Specifying Constraints in SQL
  - 6.3 – Basic Retrieval Queries in SQL
  - 6.4 – **INSERT**, **DELETE**, and **UPDATE** Statements in SQL
  - 6.5 – Additional Features of SQL

## **6.1 – SQL Data Definition and Data Types**

- SQL (Structured Query Language)
  - Considered one of the major reasons for the commercial success of relational databases
  - Appeared 1974
  - Last stable release SQL:2011
  - Core specification - Standards
  - ANSI (since 1986)
  - ISO/IEC 9075 (since 1987)

# Basic SQL

---



- **Statements for data definitions, queries, and updates**
  - DDL, DML, and VDL
  - Plus specialized **extensions** (which may be implementation specific)

# Basic SQL

---



- SQL *language* is case insensitive
  - keywords
  - namespaces
- SQL *data values* are case sensitive
- For readability, some case *style* conventions may be used
- Each statement in SQL ends with a semicolon
  - with some exemptions (e.g. USE)

# Schema and Catalog Concepts in SQL

---



- **SQL schema (In most systems, a Database)**
  - Identified by a **schema name**
  - Includes an **authorization identifier** and **descriptors** for each element
- **Schema elements include**
  - Tables
  - Constraints
  - Views
  - Domains
  - and other constructs

# SQL Terminology

---



- Terminology:
  - **Table**, **row**, and **column** used for relational model terms relation, tuple, and attribute
- **CREATE** statement
  - Main SQL command for data definition
- Much of what you'll see in the Data Definition Language is normally done with user-friendly tools like SQL Server Management Studio, etc.



# Schema and Catalog Concepts in SQL



- **CREATE SCHEMA** statement
  - `CREATE SCHEMA schema_name [auth];`
  - e.g.  
`CREATE SCHEMA COMPANY AUTHORIZATION 'Jsmith';`
- **Catalog**
  - Named collection of schemas in an *SQL environment*
- **SQL environment**
  - Installation of an SQL-compliant RDBMS on a computer system

# The CREATE TABLE Command in SQL

---



- Specify a new relation (table)
  - Provide name
  - Specify attributes and initial constraints
- Can optionally specify schema:
  - `CREATE TABLE COMPANY.EMPLOYEE . . .`
  - or
  - `CREATE TABLE EMPLOYEE . . .`

# The CREATE TABLE Command in SQL

---



- **Base tables (base relations)**
  - Relation and its tuples are actually created and stored as a file by the DBMS
- **Virtual relations**
  - Created through the **CREATE VIEW** statement

# CREATE TABLES for Company Data



```
CREATE TABLE EMPLOYEE
( Fname          VARCHAR(15)          NOT NULL,
  Minit          CHAR,
  Lname         VARCHAR(15)          NOT NULL,
  Ssn           CHAR(9)             NOT NULL,
  Bdate         DATE,
  Address       VARCHAR(30),
  Sex           CHAR,
  Salary        DECIMAL(10,2),
  Super_ssn     CHAR(9),
  Dno           INT                 NOT NULL,
  PRIMARY KEY (Ssn),
  FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn),
  FOREIGN KEY (Dno) REFERENCES DEPARTMENT(Dnumber) );

CREATE TABLE DEPARTMENT
( Dname          VARCHAR(15)          NOT NULL,
  Dnumber        INT                 NOT NULL,
  Mgr_ssn       CHAR(9)             NOT NULL,
  Mgr_start_date DATE,
  PRIMARY KEY (Dnumber),
  UNIQUE (Dname),
  FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn) );
```

**Figure 4.1**  
SQL CREATE TABLE  
data definition state-  
ments for defining the  
COMPANY schema  
from Figure 3.7.

# CREATE TABLES for Company Data



```
CREATE TABLE DEPT_LOCATIONS
( Dnumber          INT                NOT NULL,
  Dlocation        VARCHAR(15)       NOT NULL,
  PRIMARY KEY (Dnumber, Dlocation),
  FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber) );

CREATE TABLE PROJECT
( Pname          VARCHAR(15)         NOT NULL,
  Pnumber        INT                 NOT NULL,
  Plocation      VARCHAR(15),
  Dnum           INT                 NOT NULL,
  PRIMARY KEY (Pnumber),
  UNIQUE (Pname),
  FOREIGN KEY (Dnum) REFERENCES DEPARTMENT(Dnumber) );

CREATE TABLE WORKS_ON
( Essn           CHAR(9)             NOT NULL,
  Pno            INT                 NOT NULL,
  Hours         DECIMAL(3,1)         NOT NULL,
  PRIMARY KEY (Essn, Pno),
  FOREIGN KEY (Essn) REFERENCES EMPLOYEE(Ssn),
  FOREIGN KEY (Pno) REFERENCES PROJECT(Pnumber) );

CREATE TABLE DEPENDENT
( Essn           CHAR(9)             NOT NULL,
  Dependent_name VARCHAR(15)         NOT NULL,
  Sex            CHAR,
  Bdate         DATE,
  Relationship    VARCHAR(8),
  PRIMARY KEY (Essn, Dependent_name),
  FOREIGN KEY (Essn) REFERENCES EMPLOYEE(Ssn) );
```

**Figure 4.1**  
SQL CREATE TABLE  
data definition state-  
ments for defining the  
COMPANY schema  
from Figure 3.7.

## CREATE TABLE Command

---



- Does the order of table creation matter?
- Some foreign keys may cause errors
- Specified either via:
  - Circular references
  - Or because they refer to a table that has not yet been created

# Attribute Data Types and Domains in SQL

---



- Different dialects of SQL may have different types
  - Microsoft SQL Server
  - Oracle 11g, 10g, XE, etc.
  - MySQL
  - IBM DB2
  - PostgreSQL
  - SQLite

# Attribute Data Types and Domains in SQL



## ■ Basic data types

### ■ Numeric data types

- Integer numbers: INTEGER, INT, and SMALLINT
- Floating-point (real) numbers: FLOAT or REAL, and DOUBLE PRECISION

### ■ Character-string data types

- Fixed length: CHAR ( $n$ ) , CHARACTER ( $n$ )
- Varying length : VARCHAR ( $n$ ) , CHAR VARYING ( $n$ ) , CHARACTER VARYING ( $n$ )



# Attribute Data Types and Domains in SQL

---



- **Bit-string data types**
  - Fixed length: `BIT (n)`
  - Varying length: `BIT VARYING (n)`
- **Boolean data type**
  - Values of `TRUE` or `FALSE` or `NULL`
- **DATE data type (Use DateTime instead)**
  - Ten positions
  - Components are `YEAR`, `MONTH`, and `DAY` in the form `YYYY-MM-DD`

# Attribute Data Types and Domains in SQL

---



## ■ Additional data types

### ■ **Timestamp** data type (TIMESTAMP)

- Includes the DATE and TIME fields
- Plus a minimum of six positions for decimal fractions of seconds
- Optional WITH TIME ZONE qualifier

### ■ **INTERVAL** data type

- Specifies a relative value that can be used to increment or decrement an absolute value of a date, time, or timestamp

# Attribute Data Types and Domains in SQL

---



## ■ Custom Domain

- Name used with the attribute specification
- Makes it easier to change the data type for a domain that is used by numerous attributes
- Improves schema readability
- Example:
  - **CREATE DOMAIN SSN\_TYPE AS CHAR(9) ;**

## **6.2 – Specifying Constraints in SQL**

# Specifying Constraints in SQL

---



- Basic constraints:
  - Key and referential integrity constraints
  - Restrictions on attribute domains and NULLs
  - Constraints on individual tuples within a relation

# Giving Names to Constraints

---



- Keyword **CONSTRAINT**
  - Explicitly name a constraint
  - Useful for later altering

# Specifying Attribute Constraints and Attribute



- **NOT NULL**

- **NULL** is not permitted for a particular attribute

- **Default value**

- **DEFAULT** <value>

- **CHECK clause**

- **Dnumber INT NOT NULL CHECK  
(Dnumber > 0 AND Dnumber < 21) ;**

## Specifying Constraints on Tuples Using CHECK

---



- **CHECK** clauses at the end of a **CREATE TABLE** statement
  - Apply to each tuple individually
  - **CHECK (Dept\_create\_date <= Mgr\_start\_date) ;**



# DEFAULT Clause



```
CREATE TABLE EMPLOYEE
(
    Dno      INT          NOT NULL      DEFAULT 1,
    CONSTRAINT EMPPK
    PRIMARY KEY (Ssn),
    CONSTRAINT EMPSUPERFK
    FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn)
        ON DELETE SET NULL      ON UPDATE CASCADE,
    CONSTRAINT EMPDEPTFK
    FOREIGN KEY (Dno) REFERENCES DEPARTMENT(Dnumber)
        ON DELETE SET DEFAULT   ON UPDATE CASCADE);

CREATE TABLE DEPARTMENT
(
    Mgr_ssn  CHAR(9)      NOT NULL      DEFAULT '888665555',
    ...,
    CONSTRAINT DEPTPK
    PRIMARY KEY (Dnumber),
    CONSTRAINT DEPTSK
    UNIQUE (Dname),
    CONSTRAINT DEPTMGRFK
    FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn)
        ON DELETE SET DEFAULT   ON UPDATE CASCADE);

CREATE TABLE DEPT_LOCATIONS
(
    ...,
    PRIMARY KEY (Dnumber, Dlocation),
    FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber)
        ON DELETE CASCADE      ON UPDATE CASCADE);
```

**Figure 4.2**

Example illustrating how default attribute values and referential integrity triggered actions are specified in SQL.

# Specifying Key and Referential Integrity Constraints

---

## ■ PRIMARY KEY clause

- Specifies one or more attributes that make up the primary key of a relation
- **Dnumber INT PRIMARY KEY;**

## ■ UNIQUE clause

- Specifies alternate (secondary) keys
- **Dname VARCHAR(15) UNIQUE;**

# Specifying Key and Referential Integrity Constraints

---

## ■ FOREIGN KEY clause

- Default operation: reject update on violation
- Attach **referential triggered action** clause
  - Options include **SET NULL**, **CASCADE**, and **SET DEFAULT**
  - Action taken by the DBMS for **SET NULL** or **SET DEFAULT** is the same for both **ON DELETE** and **ON UPDATE**
  - **CASCADE** option suitable for “relationship” relations

## **6.3 – Basic Retrieval Queries in SQL**

# Basic Retrieval Queries in SQL

---



- **SELECT** statement
  - One basic statement for retrieving information from a database
- SQL allows a table to have two or more tuples that are identical in all their attribute values
  - Unlike relational model
  - Multi-set or bag behavior

# The Structure of Basic SQL Queries



## Basic form of the **SELECT** statement:

```
SELECT    <attribute list>  
FROM      <table list>  
WHERE     <condition>;
```

where

- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
- <table list> is a list of the relation names required to process the query.
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.

## ■ Projection attributes

- SELECT

- Attributes whose values are to be retrieved

## ■ Selection condition

- WHERE

- Boolean condition that must be true for any retrieved tuple

## ■ Logical comparison operators

- =, <, <=, >, >=, and <>

# Some Queries



**Figure 4.3**

Results of SQL queries when applied to the COMPANY database state shown in Figure 3.6. (a) Q0. (b) Q1. (c) Q2. (d) Q8. (e) Q9. (f) Q10. (g) Q1C.

(a)

| <u>Bdate</u> | <u>Address</u>          |
|--------------|-------------------------|
| 1965-01-09   | 731Fondren, Houston, TX |

(b)

| <u>Fname</u> | <u>Lname</u> | <u>Address</u>           |
|--------------|--------------|--------------------------|
| John         | Smith        | 731 Fondren, Houston, TX |
| Franklin     | Wong         | 638 Voss, Houston, TX    |
| Ramesh       | Narayan      | 975 Fire Oak, Humble, TX |
| Joyce        | English      | 5631 Rice, Houston, TX   |

**Query 0.** Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

**Q0:**     **SELECT**     Bdate, Address  
          **FROM**     EMPLOYEE  
          **WHERE**    Fname='John' **AND** Minit='B' **AND** Lname='Smith';

**Query 1.** Retrieve the name and address of all employees who work for the 'Research' department.



# Some Queries



**Figure 4.3**

Results of SQL queries when applied to the COMPANY database state shown in Figure 3.6. (a) Q0. (b) Q1. (c) Q2. (d) Q8. (e) Q9. (f) Q10. (g) Q1C.

(a)

| <u>Bdate</u> | <u>Address</u>          |
|--------------|-------------------------|
| 1965-01-09   | 731Fondren, Houston, TX |

(b)

| <u>Fname</u> | <u>Lname</u> | <u>Address</u>           |
|--------------|--------------|--------------------------|
| John         | Smith        | 731 Fondren, Houston, TX |
| Franklin     | Wong         | 638 Voss, Houston, TX    |
| Ramesh       | Narayan      | 975 Fire Oak, Humble, TX |
| Joyce        | English      | 5631 Rice, Houston, TX   |

**Query 0.** Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

**Q0:**     **SELECT**     Bdate, Address  
          **FROM**     EMPLOYEE  
          **WHERE**    Fname='John' **AND** Minit='B' **AND** Lname='Smith';

**Query 1.** Retrieve the name and address of all employees who work for the 'Research' department.

**Q1:**     **SELECT**     Fname, Lname, Address  
          **FROM**     EMPLOYEE, DEPARTMENT  
          **WHERE**    Dname='Research' **AND** Dnumber=Dno;

## ■ Query 1

- The condition  $Dnumber = Dno$  is called a **join condition**, because it combines two tuples: one from DEPARTMENT and one from EMPLOYEE, whenever the value of Dnumber in DEPARTMENT is equal to the value of Dno in EMPLOYEE.

## Some Queries



**Figure 4.3**

Results of SQL queries when applied to the COMPANY database state shown in Figure 3.6. (a) Q0. (b) Q1. (c) Q2. (d) Q8. (e) Q9. (f) Q10. (g) Q1C.

(c)

| <u>Pnumber</u> | <u>Dnum</u> | <u>Lname</u> | <u>Address</u>         | <u>Bdate</u> |
|----------------|-------------|--------------|------------------------|--------------|
| 10             | 4           | Wallace      | 291Berry, Bellaire, TX | 1941-06-20   |
| 30             | 4           | Wallace      | 291Berry, Bellaire, TX | 1941-06-20   |

**Query 2.** For every project located in ‘Stafford’, list the project number, the controlling department number, and the department manager’s last name, address, and birth date.

## Some Queries



**Figure 4.3**

Results of SQL queries when applied to the COMPANY database state shown in Figure 3.6. (a) Q0. (b) Q1. (c) Q2. (d) Q8. (e) Q9. (f) Q10. (g) Q1C.

(c)

| <u>Pnumber</u> | <u>Dnum</u> | <u>Lname</u> | <u>Address</u>         | <u>Bdate</u> |
|----------------|-------------|--------------|------------------------|--------------|
| 10             | 4           | Wallace      | 291Berry, Bellaire, TX | 1941-06-20   |
| 30             | 4           | Wallace      | 291Berry, Bellaire, TX | 1941-06-20   |

**Query 2.** For every project located in ‘Stafford’, list the project number, the controlling department number, and the department manager’s last name, address, and birth date.

**Q2:**     **SELECT**     Pnumber, Dnum, Lname, Address, Bdate  
          **FROM**     PROJECT, DEPARTMENT, EMPLOYEE  
          **WHERE**    Dnum=Dnumber **AND** Mgr\_ssn=Ssn **AND**  
                    Plocation=‘Stafford’;

## Ambiguous Attribute Names



- Same name can be used for two (or more) attributes
  - As long as the attributes are in different relations
  - Must **qualify** the attribute name with the relation name to prevent ambiguity

```
Q1A:  SELECT  Fname, EMPLOYEE.Name, Address
      FROM    EMPLOYEE, DEPARTMENT
      WHERE   DEPARTMENT.Name='Research' AND
              DEPARTMENT.Dnumber=EMPLOYEE.Dnumber;
```

# Aliasing, Renaming, and Tuple Variables



- The SQL **AS** keyword
- Attribute or Table Aliases
- Examples
  - SELECT *attribute* AS *attr\_alias\_name*
  - FROM *table\_name* AS *table\_alias\_name*
  - FROM (*subquery*) AS *virtual\_table\_name*
- CAVEAT
  - SQL syntax allows AS to be implied
  - A missing comma in a SELECT clause can be parsed as an implied AS

# Unspecified WHERE Clause



- Missing WHERE clause
  - Indicates no condition on tuple selection
- CROSS PRODUCT
  - All possible tuple combinations

Queries 9 and 10. Select all EMPLOYEE Ssns (Q9) and all combinations of EMPLOYEE Ssn and DEPARTMENT Dname (Q10) in the database.

Q9:     SELECT     Ssn  
          FROM     EMPLOYEE;

Q10:    SELECT     Ssn, Dname  
          FROM     EMPLOYEE, DEPARTMENT;

# Unspecified WHERE Clause and Use of the Asterisk



## ■ Specify an asterisk (\*)

- Retrieve all the attribute values of the selected tuples

|       |        |                                   |
|-------|--------|-----------------------------------|
| Q1C:  | SELECT | *                                 |
|       | FROM   | EMPLOYEE                          |
|       | WHERE  | Dno=5;                            |
| Q1D:  | SELECT | *                                 |
|       | FROM   | EMPLOYEE, DEPARTMENT              |
|       | WHERE  | Dname='Research' AND Dno=Dnumber; |
| Q10A: | SELECT | *                                 |
|       | FROM   | EMPLOYEE, DEPARTMENT;             |



## Tables as Sets in SQL



- SQL does not automatically eliminate duplicate tuples in query results
- Use the keyword **DISTINCT** in the **SELECT** clause
  - Only distinct tuples should remain in the result

**Query 11.** Retrieve the salary of every employee (Q11) and all distinct salary values (Q11A).

**Q11:**    **SELECT**    **ALL** Salary  
          **FROM**     EMPLOYEE;

**Q11A:**  **SELECT**    **DISTINCT** Salary  
          **FROM**     EMPLOYEE;

# Tables as Sets in SQL



## ■ Set operations

- **UNION, INTERSECT, EXCEPT** (Set difference)
- Corresponding multiset operations: **UNION ALL, EXCEPT ALL, INTERSECT ALL**
- Macros around binary set operators

**Query 4.** Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

```
Q4A: (SELECT DISTINCT Pnumber
      FROM PROJECT, DEPARTMENT, EMPLOYEE
      WHERE Dnum=Dnumber AND Mgr_ssn=Ssn
            AND Lname='Smith' )

      UNION

      ( SELECT DISTINCT Pnumber
        FROM PROJECT, WORKS_ON, EMPLOYEE
        WHERE Pnumber=Pno AND Essn=Ssn
              AND Lname='Smith' );
```

# Substring Pattern Matching and Arithmetic Operators

---



- **LIKE** comparison operator
  - Used for string **pattern matching**
  - % replaces an arbitrary number of zero or more characters
  - underscore (\_) replaces a single character
- Standard arithmetic operators:
  - Addition (+), subtraction (−), multiplication (\*), and division (/)
- **BETWEEN** comparison operator

## Ordering of Query Results



- SQL results are inherently not ordered. To change this, use **ORDER BY** clause
  - Keyword **DESC** to see result in a descending order of values
  - Keyword **ASC** to specify ascending order explicitly
  - **ORDER BY D.Dname DESC, E.Lname ASC, E.Fname ASC**

# Discussion and Summary of Basic SQL Retrieval Queries

---



```
SELECT    <attribute list>
FROM      <table list>
[ WHERE   <condition> ]
[ ORDER BY <attribute list> ];
```

## **6.4 – INSERT, DELETE, and UPDATE Statements in SQL**

# Modifying the Database Data

---



- Three commands used to modify the database:
  - **INSERT**
  - **DELETE**
  - **UPDATE**

# The INSERT Command



- Specify the relation name and a list of values for the tuple

```
U1:  INSERT INTO  EMPLOYEE
      VALUES      ( 'Richard', 'K', 'Marini', '653298653', '1962-12-30', '98
                    Oak Forest, Katy, TX', 'M', 37000, '653298653', 4 );
```



# The INSERT Command



- Specify the relation name and a list of values for the tuple

```
U3B:  INSERT INTO  WORKS_ON_INFO ( Emp_name, Proj_name,  
                Hours_per_week )  
      SELECT      E.Lname, P.Pname, W.Hours  
      FROM        PROJECT P, WORKS_ON W, EMPLOYEE E  
      WHERE       P.Pnumber=W.Pno AND W.Essn=E.Ssn;
```

# The CREATE TABLE Command



- Define auto-increment at table create time

```
CREATE TABLE Persons (  
    ID int NOT NULL AUTO_INCREMENT PRIMARY KEY,  
    LastName VARCHAR(255) NOT NULL,  
    FirstName VARCHAR(255),  
    Address VARCHAR(255),  
    City VARCHAR(255)  
);
```

- Define/Add auto-increment to an existing table

```
ALTER TABLE Persons AUTO_INCREMENT=100
```

# The DELETE Command



- Removes tuples from a relation
  - Includes a WHERE clause to select the tuples to be deleted

|      |             |                  |
|------|-------------|------------------|
| U4A: | DELETE FROM | EMPLOYEE         |
|      | WHERE       | Lname='Brown';   |
| U4B: | DELETE FROM | EMPLOYEE         |
|      | WHERE       | Ssn='123456789'; |
| U4C: | DELETE FROM | EMPLOYEE         |
|      | WHERE       | Dno=5;           |
| U4D: | DELETE FROM | EMPLOYEE;        |

## The UPDATE Command



- Modify attribute values of one or more selected tuples
- Additional **SET** clause in the UPDATE command
  - Specifies attributes to be modified and new values

```
U5:  UPDATE  PROJECT
      SET     Plocation = 'Bellaire', Dnum = 5
      WHERE   Pnumber=10;
```

## Additional Features of SQL

---



- Techniques for specifying complex retrieval queries
- Writing programs in various programming languages that include SQL statements
- Set of commands for specifying physical database design parameters, file structures for relations, and access paths
- Transaction control commands

## Additional Features of SQL

---



- Specifying the granting and revoking of privileges to users
- Constructs for creating triggers
- Enhanced relational systems known as object-relational
- New technologies such as XML and OLAP

# Summary



## ■ SQL

- Comprehensive language
- Data definition, queries, updates, constraint specification, and view definition

## ■ Covered in Chapter 6:

- Data definition commands for creating tables
- Commands for constraint specification
- Simple retrieval queries
- Database modification commands