

Gaussian Naive Bayesian

(Taken mostly from the book by Duda, Hart, and Stork.)

Gaussian density/distribution

The Gaussian density function of n -dimensional vectors is:

$$g(x; \mu, C) = \frac{1}{(\sqrt{2\pi})^n |C|^{1/2}} e^{-\frac{1}{2}(x-\mu)^T C^{-1}(x-\mu)}$$

Here μ is the distribution mean and C is the Covariance matrix. The value $|C|$ is the determinant of the matrix C . The parameters μ, C can be estimated from the data by:

$$\mu = \frac{\sum_{i=1}^m x_i}{m}, \quad C = \frac{\sum_i (x_i - \mu)(x_i - \mu)^T}{m} \quad \text{or} \quad C = \frac{\sum_i (x_i - \mu)(x_i - \mu)^T}{m - 1}$$

$$\text{The special } n=1 \text{ case is:} \quad g(x; \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2\sigma^2}(x-\mu)^2}$$

The parameters μ, σ can be estimated from the data by:

$$\mu = \frac{\sum_{i=1}^m x_i}{m}, \quad C = \frac{\sum_i (x_i - \mu)^2}{m} \quad \text{or} \quad C = \frac{\sum_i (x_i - \mu)^2}{m - 1}, \quad \sigma = \sqrt{C}$$

Discriminant functions for Gaussian density

Suppose the positive examples have Gaussian density with parameters μ_1, C_1 , and the negative examples have Gaussian density with parameters μ_2, C_2 . When should we decide that x is positive and not negative? Call h_1 the hypothesis that x is positive, and h_2 the hypothesis that x is negative. We should decide that x is positive if: $P(h_1|x) > P(h_2|x)$. Observe that $P(x|h_1) = g(x; \mu_1, C_1)$, and similarly $P(x|h_2) = g(x; \mu_2, C_2)$. Therefore:

$$\begin{aligned} \text{ML classifier: Classify as positive if} & \quad g(x; \mu_1, C_1) > g(x; \mu_2, C_2) \\ \text{MAP classifier: Classify as positive if} & \quad g(x; \mu_1, C_1)P(h_1) > g(x; \mu_2, C_2)P(h_2) \end{aligned}$$

Concentrating on the MAP classifier the condition $g(x; \mu_1, C_1)P(h_1) > g(x; \mu_2, C_2)P(h_2)$ gives the following after taking logarithm:

$$\begin{aligned} & -\frac{1}{2}(x - \mu_1)^T C_1^{-1}(x - \mu_1) - \frac{n \ln 2\pi}{2} - \frac{\ln |C_1|}{2} + \ln P(h_1) \\ & > -\frac{1}{2}(x - \mu_2)^T C_2^{-1}(x - \mu_2) - \frac{n \ln 2\pi}{2} - \frac{\ln |C_2|}{2} + \ln P(h_2) \end{aligned}$$

This simplifies to:

$$(x - \mu_2)^T C_2^{-1}(x - \mu_2) - (x - \mu_1)^T C_1^{-1}(x - \mu_1) + \ln |C_2| - \ln |C_1| + 2(\ln P(h_1) - \ln P(h_2)) > 0$$

This is a quadratic discriminant function. The main problem with this approach is that the Covariance matrices have $O(n^2)$ variables that need to be determined, and reliable estimation may require a lot of training data. Instead, in many situations this model is used with additional simplifying assumptions that reduce the quadratic expression to a linear expression.

Case 1: $C_1 = C_2 = \sigma^2 I$

Here the discriminant function is:

$$|x - \mu_2|^2 - |x - \mu_1|^2 + 2\sigma^2(\ln P(h_1) - \ln P(h_2)) > 0$$

This can be further simplified to:

$$\begin{aligned} d(x) &= w^T x + b > 0 \quad \text{where:} \\ w &= (\mu_1 - \mu_2) \\ b &= \frac{1}{2}(|\mu_2|^2 - |\mu_1|^2) + \sigma^2(\ln P(h_1) - \ln P(h_2)) \end{aligned}$$

Typically only the value of w is computed. The values of $w^T x_i$ are calculated for all i , and the value of b is determined using an ad-hoc technique for computing a threshold.

Case 2: $C_1 = C_2 = C$

In this case it can be shown that the discriminant function is:

$$\begin{aligned} d(x) &= w^T x + b > 0 \quad \text{where:} \\ w &= C^{-1}(\mu_1 - \mu_2) \\ b &= \frac{1}{2}(\mu_2^T C^{-1} \mu_2 - \mu_1^T C^{-1} \mu_1) + (\ln P(h_1) - \ln P(h_2)) \end{aligned}$$

Typically only the value of w is computed. The values of $w^T x_i$ are calculated for all i , and the value of b is determined using an ad-hoc technique for computing a threshold.

Case 3: arbitrary C_1, C_2

Here we can get an arbitrary quadratic function as the discriminant.

$$\begin{aligned} d(x) &= (x - \mu_2)^T C_2^{-1} (x - \mu_2) - (x - \mu_1)^T C_1^{-1} (x - \mu_1) + \ln |C_2| - \ln |C_1| + 2(\ln P(h_1) - \ln P(h_2)) \\ &= (x - \mu_2)^T C_2^{-1} (x - \mu_2) - (x - \mu_1)^T C_1^{-1} (x - \mu_1) + b > 0 \end{aligned}$$

Typically the value of b is determined using an ad-hoc technique for computing a threshold.

Example

x_1	x_2	y
2	6	+
3	8	+
4	6	+
3	4	+
1	-2	-
3	0	-
5	-2	-
3	-4	-

.	.	+	.	.
.
.	+	.	+	.
.
.	.	+	.	.
.
.
.	.	-	.	.
.
-	.	.	.	-
.
.	.	-	.	.

Case 1:

$$\mu_1 = \begin{pmatrix} 3 \\ 6 \end{pmatrix}, \quad \mu_2 = \begin{pmatrix} 3 \\ -2 \end{pmatrix}$$

$$w = \begin{pmatrix} 0 \\ 8 \end{pmatrix}, \quad b = -16$$

$$d(x) = 8x_2 - 16, \quad \text{or} \quad d(x) = x_2 - 2$$

The value of b was computed by considering the sorted values of $w^T x$. They are:

$-4 \cdot 8$	$-2 \cdot 8$	$-2 \cdot 8$	$0 \cdot 8$	$4 \cdot 8$	$6 \cdot 8$	$6 \cdot 8$	$8 \cdot 8$
-	-	-	-	+	+	+	+

The value of b that gives the smallest error and maximizes the margins.

Case 2:

$$\mu = \begin{pmatrix} 3 \\ 2 \end{pmatrix}$$

$$C = \frac{1}{8} \begin{pmatrix} 10 & 0 \\ 0 & 144 \end{pmatrix} = \begin{pmatrix} 10/8 & 0 \\ 0 & 18 \end{pmatrix}$$

To compute w solve the following linear system:

$$Cw = \mu_1 - \mu_2$$

This gives:

$$w = \begin{pmatrix} 0 \\ 4/9 \end{pmatrix}$$

$$d(x) = \frac{4}{9}x_2 + b \quad \text{or} \quad d(x) = x_2 + b = x_2 - 2$$

The value of b was computed in the same way as in Case 1. The result is the same as in Case 1: $b = -2$.

Case 3:

$$C_1 = \begin{pmatrix} 1/2 & 0 \\ 0 & 2 \end{pmatrix}, \quad C_1^{-1} = \begin{pmatrix} 2 & 0 \\ 0 & 1/2 \end{pmatrix}$$

$$C_2 = \begin{pmatrix} 2 & 0 \\ 0 & 2 \end{pmatrix}, \quad C_2^{-1} = \begin{pmatrix} 1/2 & 0 \\ 0 & 1/2 \end{pmatrix}$$

We have:

$$(x - \mu_1)^T C_1^{-1} (x - \mu_1) = 2(x_1 - 3)^2 + \frac{1}{2}(x_2 - 6)^2$$

$$(x - \mu_2)^T C_2^{-1} (x - \mu_2) = \frac{1}{2}(x_1 - 3)^2 + \frac{1}{2}(x_2 + 2)^2$$

This gives the following expression for the discriminant:

$$d(x) = \frac{1}{2}(x_1 - 3)^2 + \frac{1}{2}(x_2 + 2)^2 - 2(x_1 - 3)^2 - \frac{1}{2}(x_2 - 6)^2 + b = d_q(x) + b$$

Simplifying $d_q(x)$ we get:

$$d_q(x_1, x_2) = -1.5x_1^2 + 9x_1 + 8x_2 - 29.5$$

To compute b we consider the sorted values of $d_q(x)$:

$d_q(3, -4)$ = -48	$d_q(5, -2)$ = -38	$d_q(1, -2)$ = -38	$d_q(3, 0)$ = -16	$d_q(3, 4)$ = 16	$d_q(4, 6)$ = 30.5	$d_q(2, 6)$ = 30.5	$d_q(3, 8)$ = 48
-	-	-	-	+	+	+	+

The value of b that gives the smallest error and maximizes the margins is $b = 0$. Observe that this result is different from cases 1,2. For example, in Case 3 we classify an example as negative whenever x_1 is large positive or large negative.