

Clustering data into k clusters

Given n points in \mathbb{R}^n split them into “similar” groups. This problem is tough even if the number of groups is known, which is the only case we consider.

Input: m feature vectors x_1, \dots, x_m , and the value of k .

Output: Partitioning of the data into k classes.

There are several ways to describe the partitions. One is an index function $c(i)$:

$$c(i) = j \quad \Leftrightarrow \quad x_i \text{ belongs to class } j$$

Another one is a membership function $c(i, j)$:

$$c(i, j) = \begin{cases} 1 & x_i \text{ belongs to class } j \\ 0 & \text{otherwise} \end{cases}$$

A probabilistic, or fuzzy version of the membership function is:

$$c(i, j) = p_{i,j}$$

subject to the conditions: $p_{ij} \geq 0$, and for each i , $\sum_j p_{ij} = 1$. Here the value of p_{ij} is the probability, or the likelihood, that x_i belongs to class j .

K-means clustering

The output of K-means is sometimes called *hard partition*. Define the quantization error by:

$$E = \sum_{j=1}^k \sum_{c(i)=j} |x_i - u_j|^2 = \sum_{i=1}^m |x_i - u_{c(i)}|^2$$

Keeping $c(i)$ unchanged, the value of u_j that minimizes E is: can be computed by taking the derivative of E with respect to u_j and equating to 0. Observe that only the terms $|x_i - u_j|^2$ depend on u_j , and

$$\frac{\partial}{\partial u_j} |x_i - u_j|^2 = \frac{\partial}{\partial u_j} (|x_i|^2 - 2x_i^T u_j + |u_j|^2) = 2u_j - 2x_i$$

Therefore:

$$\frac{\partial E}{\partial u_j} = \sum_{c(i)=j} (u_j - x_i) = m_j u_j - \sum_{c(i)=j} x_i$$

where $m_j = \sum_{c(i)=j} 1$ is the number of points in class j . This gives:

$$\text{for all } j, \quad u_j = \frac{\sum_{c(i)=j} x_i}{m_j} \tag{1}$$

Keeping u_j unchanged, the value of $c(i)$ that minimizes E is:

$$\text{for all } i, \quad c(i) = \arg \min_j |x_i - u_j|^2 \tag{2}$$

Typical algorithm:

0. Start with a guess for u_1, \dots, u_k , or a guess for the $c(i)$.
1. Iterate until convergence the equations (1),(2).

It is easy to verify that each iteration reduced the error E , which proves convergence. Unfortunately the algorithm “almost always” gets stuck in a local minimum. The standard approach is to run it many times with different initial estimates, and select the partitioning that gives the smallest value of E .

Lloyd’s algorithm

Lloyd described a specific formulation of randomizing the initial step of the algorithm. In Lloyd’s version Step 0 is implemented as follows:

0. Select the initial u_i uniformly at random from x_1, \dots, x_m .

Kmeans++

Here we spend more time on the initial step, the initial selection of u_1, \dots, u_k .

- 0.1. Select u_1 uniformly at random from x_1, \dots, x_m .
- 0.2. Assuming that u_1, \dots, u_t are already selected, with $t < k$. Use the following procedure to select u_{t+1} :

- 0.2.1 for each x_i compute the value d_i as:

$$d_i = \min_j |x_i - u_j|^2$$

- 0.2.2 Compute the probability values:

$$Z = \sum_i d_i, \quad p_i = \frac{d_i}{Z}$$

- 0.2.3 Select one of the x_i at random according to the probabilities p_i . Set it as u_{t+1} .