

```

1  # Learning with a linear model
2
3  import numpy as np
4  import tensorflow as tf
5
6  # suppress warnings
7  old_v = tf.logging.get_verbosity()
8  tf.logging.set_verbosity(tf.logging.ERROR)
9
10 from tensorflow.examples.tutorials.mnist import input_data
11 mnist = input_data.read_data_sets('MNIST_data', one_hot=True)
12 tf.logging.set_verbosity(old_v)
13
14 # MNIST_data is a collection of 2D gray level images.
15 # Each image is a picture of a digit from 0..9
16 # Each image is of size 28 x 28 pixels
17
18 sess = tf.InteractiveSession()
19
20 # xi is an image of size n. yi is the N labels of the image
21 # X is mxn. Row xi of X is an image
22 # Y is mxN. Row yi of Y is the labels of xi
23 X = tf.placeholder(tf.float32, shape=[None, 784])
24 Y = tf.placeholder(tf.float32, shape=[None, 10])
25
26 # consider the linear model:  $W^T xi + b = yi$ . Here W is nxN, b is Nx1.
27 # In matrix form:  $X W + B = Y$ , where B is mxN.  $B = 1 b^T$ 
28
29 W = tf.Variable(tf.zeros([784,10]))
30 b = tf.Variable(tf.zeros([10]))
31
32 sess.run(tf.global_variables_initializer())
33
34 # the linear regression model: (b is one row. It will be replicated as needed.)
35
36 predicted_Y = tf.matmul(X,W) + b
37
38 # means squared error loss is very bad
39 # loss = tf.losses.mean_squared_error(labels = Y, predictions = predicted_Y)
40 loss = tf.nn.softmax_cross_entropy_with_logits_v2(labels=Y, logits=predicted_Y)
41 train_step = tf.train.GradientDescentOptimizer(0.5).minimize(loss)
42
43 for i in range(1000):
44     batch = mnist.train.next_batch(100)
45     train_step.run(feed_dict={X: batch[0], Y: batch[1]})
46
47 correct_prediction = tf.equal(tf.argmax(predicted_Y,1), tf.argmax(Y,1))
48 accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
49
50 print(accuracy.eval(feed_dict={X: mnist.test.images, Y: mnist.test.labels}))

```