

The Nearest-Neighbor and the k -Nearest-Neighbor algorithms

In their simplest form, these algorithms do not perform any computation during training. The computation is performed only when a test example is presented. Therefore, they are described with input that contains the training examples and one test example. The expensive step in these algorithms is the computation of nearest-neighbors. Efficient implementations exist with sophisticated data structures for efficient computation of nearest-neighbors. These are not discussed here.

Input: m training examples, given as the pairs (x_i, y_i) , where x_i is an n -dimensional feature vector and y_i is its label. A test example x .

Output: y , the computed label of x .

The Nearest-Neighbor algorithm

a. Determine x_i nearest to x . It minimizes the distance to x according to a pre-defined norm.

$$\text{distance}(x_i, x) = |x_i - x| \quad (1)$$

b. Return $y = y_i$.

The most commonly used norm in (1) is the Euclidean norm:

$$|x_i - x|^2 = \sum_{j=1}^n (x_i(j) - x(j))^2$$

There are multiple approaches for handling the case in which there is more than one training example nearest to x .

The k -Nearest-Neighbor algorithm

a. Determine x_{i1}, \dots, x_{ik} , the k training examples nearest to x according to a pre-defined norm.

b. Let y_{i1}, \dots, y_{ik} be the labels of the k nearest neighbors. Choose y as the label that appears most frequently among y_{i1}, \dots, y_{ik} .

There are multiple approaches for handling the case in which no label has a clear majority in b.

The value of k

In many practical problems k -NN with $k > 1$ performs better than the simple 1-NN. The most effective method of estimating a useful value of k is the technique of cross validation.

Example

Training data:	i	x_i	y_i
	1	(1,1)	A
	2	(1,2)	A
	3	(2,2)	B
	4	(2,3)	B

To classify the test example (3,2) according to k -NN we need to compute its distances to the 4 examples. The square of the Euclidean distances is shown in the following table:

i	x_i	y_i	$ x_i - (3, 2) ^2$
1	(1,1)	A	5
2	(1,2)	A	4
3	(2,2)	B	1
4	(2,3)	B	2

Using 1-NN the nearest example has the index $i = 3$, and the label is: $y = y_3 = B$.

Using 3-NN the 3 nearest examples are with indexes $i = 3, 4, 2$. Two have label B and one label A , so that the computed label is $y = B$.